

The BalANS protocol

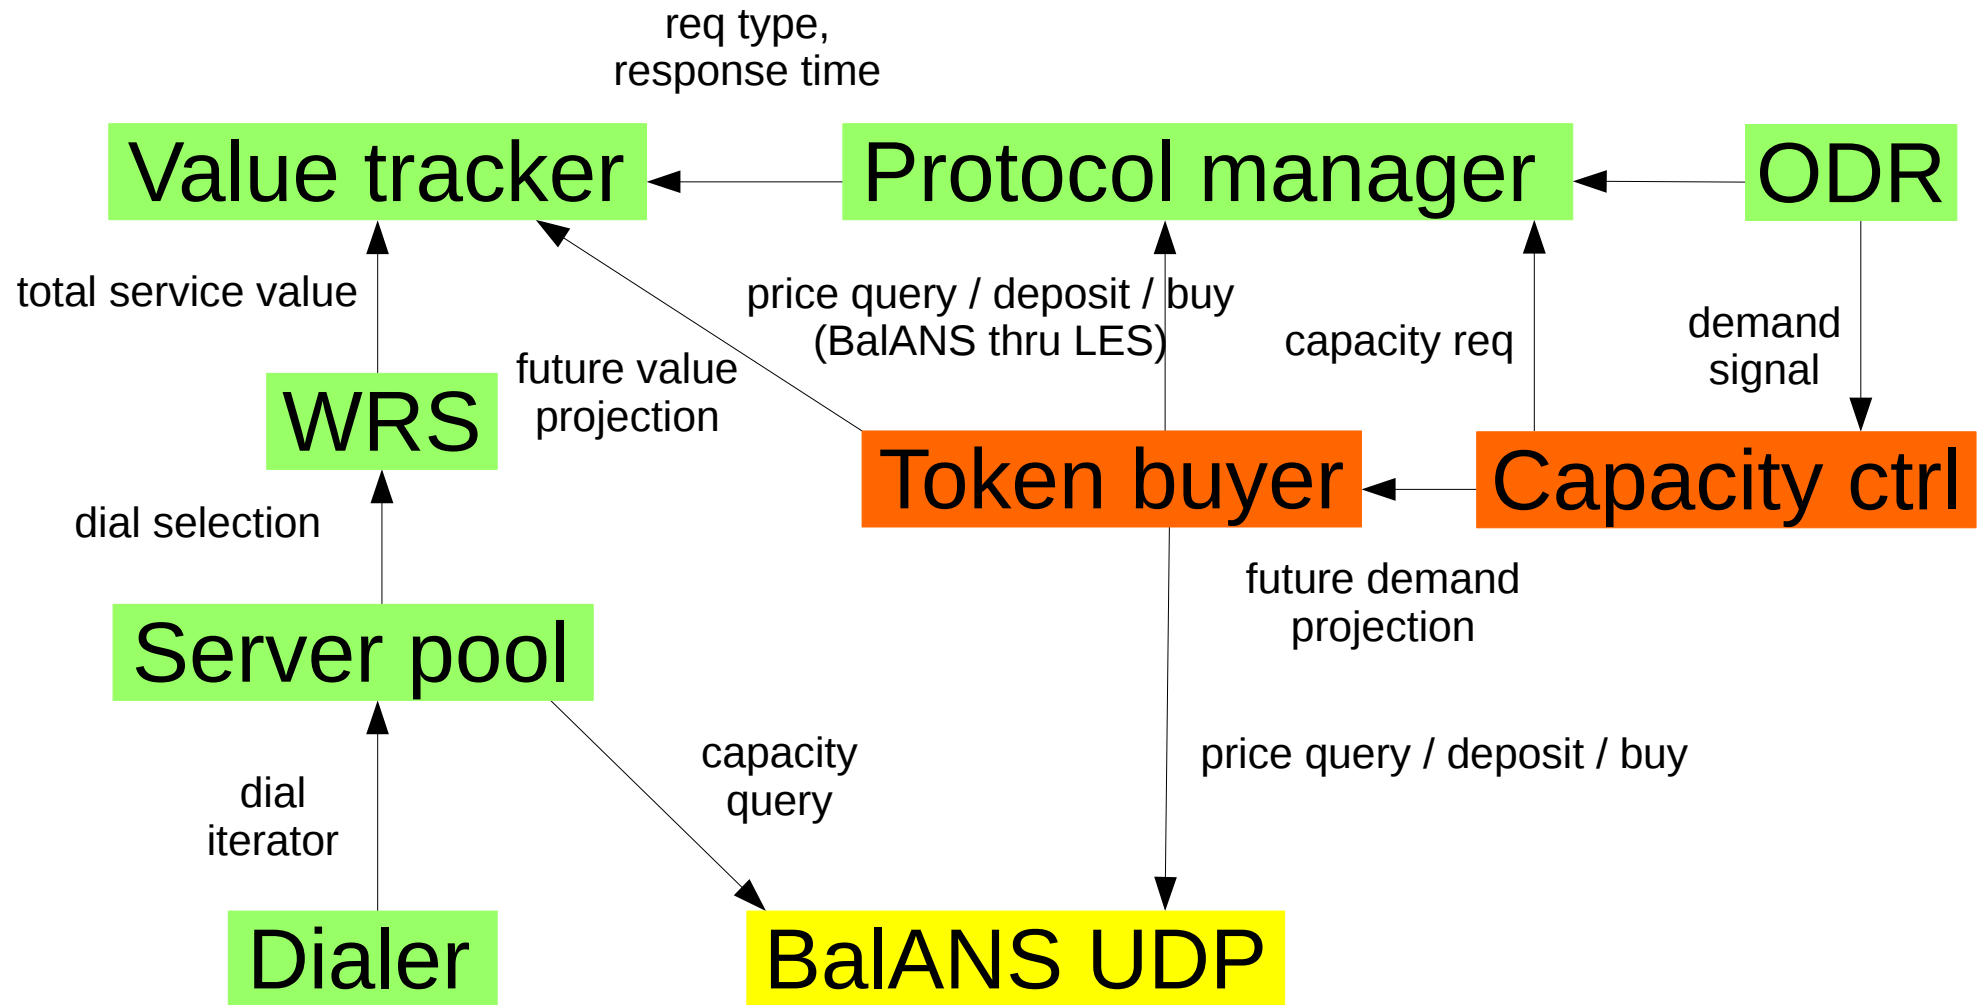
“Balanced Asymmetric Network Scaling”

- implementation and embedding into LES (protocol upgrade, code refactoring)
- decentralized economic/trust model
- smart client strategy
- future plans: more advanced network topologies (sharding, caching, relay nodes)

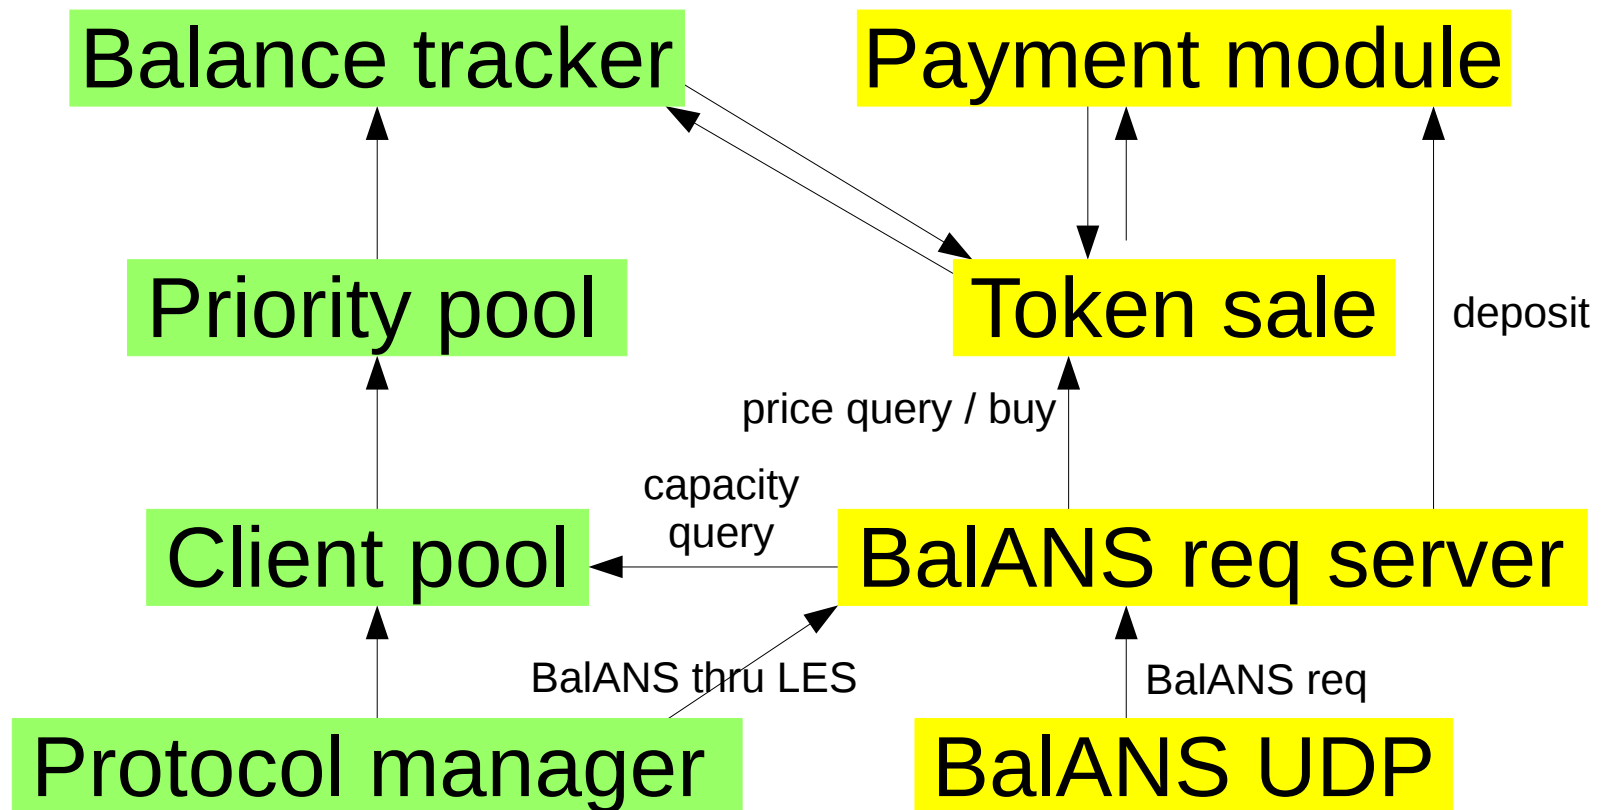
Protocol upgrades

- BalANS protocol (available through both LES/4 and discv5 UDP)
 - capacity query
 - price query
 - deposit
 - exchange
- LES/4 enhancements
 - serve BalANS requests
 - capacity request/update
 - request/reply meta fields
 - reply optionally includes real request cost and token balance (required for paid connections)

Integration and refactoring (client side)



Integration and refactoring (server side)




Economic model

“Free market planned economy”

- policies for the short term, market forces for the long term
 - adjustable policy parameters subject to market negotiation
 - two-level scalable priority model
 - capacity assigned to each client, total capacity limited
 - capacity guarantees requests being served
 - pre-purchased tokens guarantee capacity being granted
- “adjusted bonding curve” token sale mechanism
 - token utility is stable, token amount is limited
- flexible pricing
 - price tables based on worst case difficulty estimates
 - server rewards easier than worst case requests with discounts


Server policy and token sale

Max capacity/balance ratio
(enforces max total capacity)

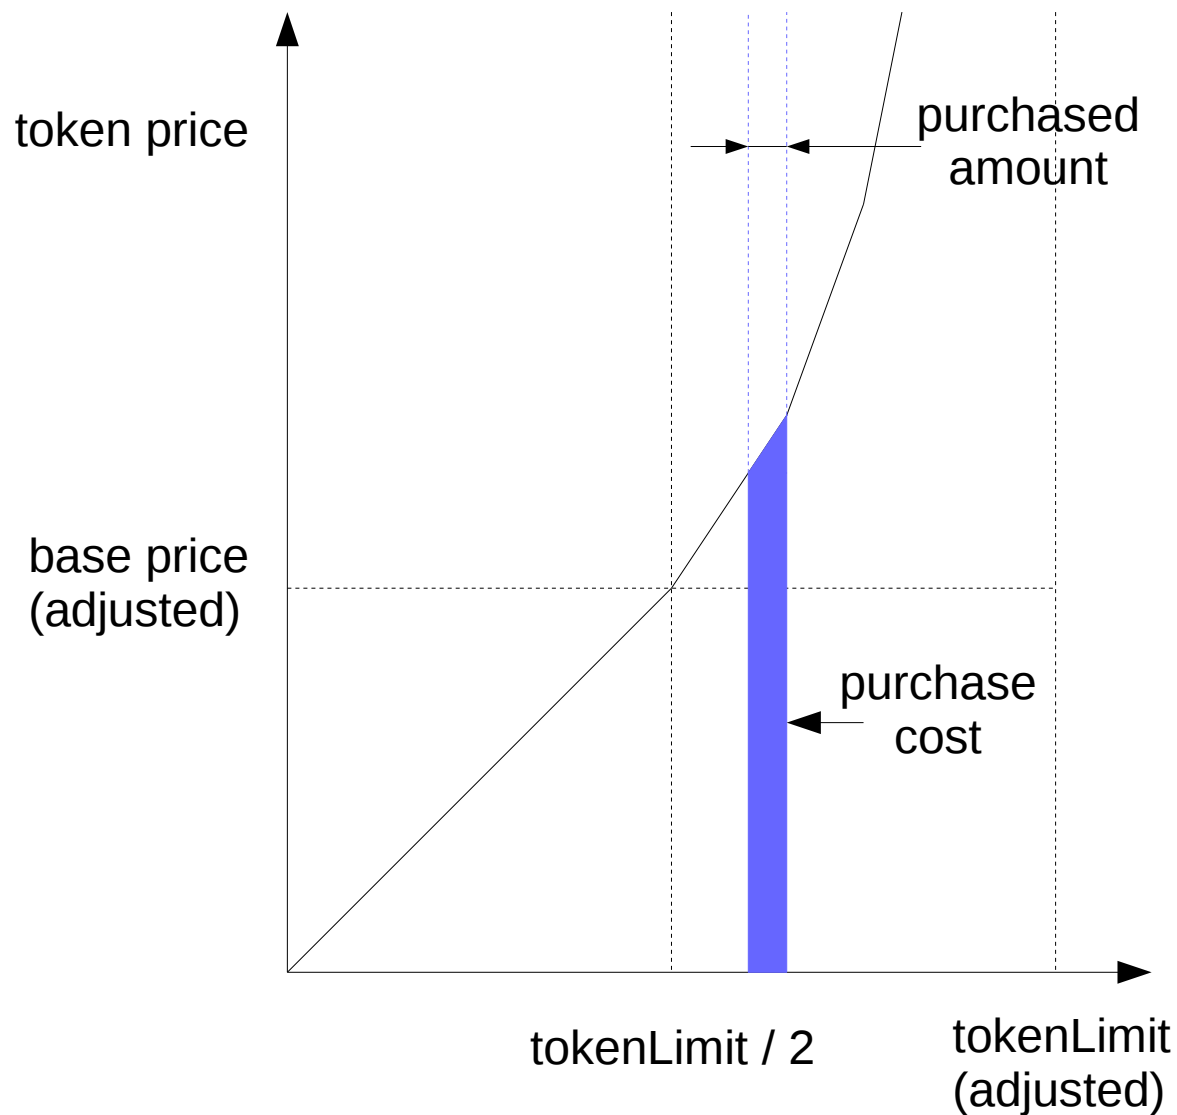


Token balance

Max req rate/capacity ratio
(enforces global req rate limit)



Capacity




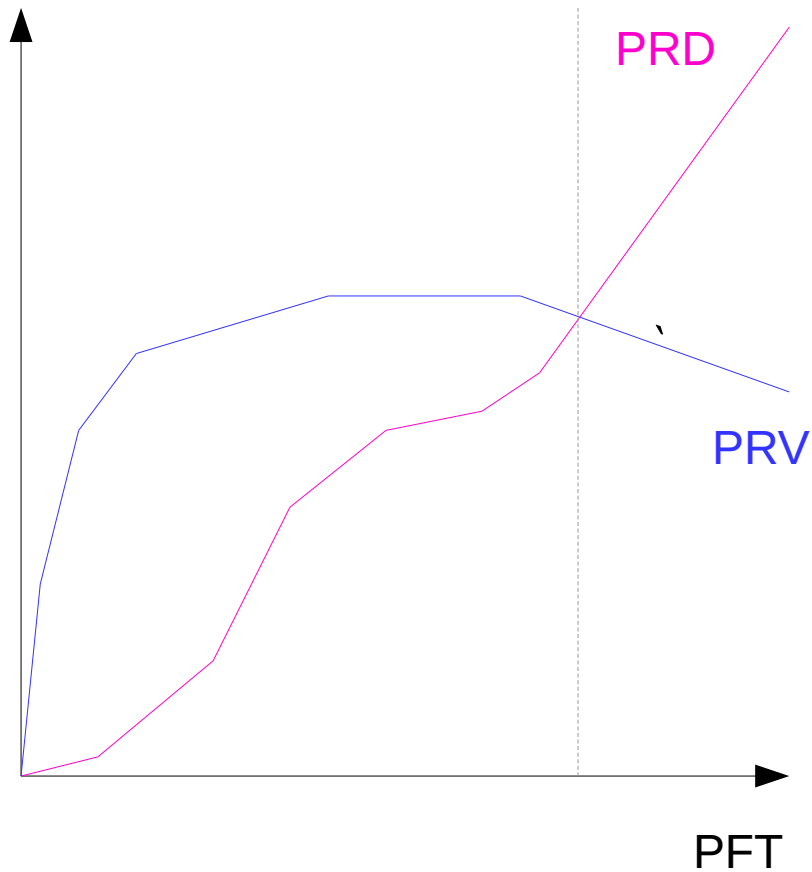
Decentralized trust model

- caveat emptor (nothing is enforceable)
- mutual strategy commitment builds trust
 - trust limits grow exponentially if both parties cooperate
 - kickstarted by free service
- smart clients, dumb servers
 - servers have a simple priority policy
 - clients build a probabilistic model of the known servers based on statistical data in order to optimize their decisions
 - this model predicts service value in the near future based on server policy parameters
 - service as a commodity (instead of a complex enterprise)
 - easy to provide, easy to access, easy to marketize
 - balances information asymmetry, more power stays with the smallest players
 - reduces entry barrier for server operation (risk is a bigger centralization factor than technological economies of scale)



Client market strategy

- a dynamic balance between allocated future time and money spent
- projectedTokenValue, $PTV(dt)$: the amount of quality corrected request value expected from all servers if we spend all tokens in the next dt time
 - based on token balance, announced request and capacity prices, token expiration speed and capacity-dependent performance correction factors based on past statistics
 - $PTV_exp(dt)$ is the amount of request value lost to token expiration in the above scenario
- projectedTokenDemand, $PTD(dt)$: a rough estimate of future request value demand (simply based on the last dt time period)
- projectedFutureTime (PFT) is where $PTV(PFT) = PTD(PFT)$
- PaymentBuffer (PB) is charged with MaxPaymentRate (MPR, currency/time) if **$PFT < FutureTarget \ \&\& \ PTV_exp(PFT) / PTV(PFT) < MaxLossRate$**
- TokenBuyer makes price and capacity queries, calculates PFT' and PB' for different possible token purchase amounts
- A purchase is made if **$PB' * PFT' > PB * PFT$**

Client market strategy



Service quality modeling



Advanced topologies

- flat topology
 - all servers provide the same service (access the entire chain)
- layered topology
 - “service map” allows advertising and looking for access to subsets of the whole dataset
 - some servers provide quick access to a subset of all data (full node on a shard)
 - other servers do request relaying and data caching of data belonging to the entire dataset or a larger subset (light client on many shards)
- “holistic” topology
 - all servers provide access to all data, with different delays on different subsets (at least a light client on every shard)
 - service map adds a “distance” dimension (proportional to maximal guaranteed delay)
 - recommendation service (gives the ENR of a proven efficient node providing a lower distance access to a given subset)
 - practically an advanced, incentivized discovery service based on real performance